

Provenance-Verifiable Privacy Cloaking for Face Images: Safety-Region-Guided Neural Watermarking with JPEG-Robust Detection

Author Name^{a,*}

^a*Department, University, City, Country*

Abstract

Privacy-cloaking techniques such as Fawkes protect face images from unauthorized face recognition (FR) by adding imperceptible perturbations; however, cloaked images carry no verifiable provenance, allowing anyone to republish a victim’s cloaked face without accountability. This paper studies *provenance-verifiable privacy cloaking*: embedding a decodable neural watermark into cloaked faces so that FR evasion is preserved, the watermark remains detectable, and detection is robust to JPEG compression. We first reveal a fundamental conflict: Fawkes cloaking reduces WAM watermark decoding from 99.8% bit accuracy (98% TPR@95%) on clean faces to 85.2% (6%) on cloaked faces, quantifying, for the first time, how adversarial cloaking damages neural watermarks. We then propose two complementary mechanisms. (1) *Safety-region guidance* (SRG) computes an FR-sensitivity map from surrogate gradients and embeds watermarks preferentially in FR-insensitive regions, raising cloaking retention under strong watermarking from 36.1% to 42.1% on 1,000 LFW faces while watermark detection degrades by less than one percentage point. (2) *Strength-calibrated embedding with BER-based detection* increases watermark detectability from 6% to 57.9% (TPR@95%) and to 99.2%/99.6%/93.7% under $\text{BER} \leq 0.25$ detection for clean/JPEG80/JPEG50, at 0% false-positive rate. Under stronger cloaking (Fawkes mid, 85% protection success rate), watermarking coexists losslessly: retention is unchanged and detection reaches 95–100%. Finally, we show that sequential embedding is competitive with joint optimization

*Corresponding author

Email address: `author@example.com` (Author Name)



Figure 1: Provenance-verifiable privacy cloaking. Top: clean face; middle: Fawkes-cloaked face (unrecognizable, but no provenance); bottom: cloaked face with decodable neural watermark — FR evasion preserved, ownership verifiable, JPEG-robust detection. (Placeholder: to be generated.)

(watermark-aware refinement), directly informing the open question raised by EIAW (BDMA 2026) and ARFP (arXiv:2605.01217). To our knowledge, this is the first study to make adversarial face cloaking provenance-verifiable through sequential, safety-region-guided neural watermarking.

Keywords: Face recognition, privacy cloaking, neural watermarking, JPEG robustness, provenance

1. Introduction

Face recognition (FR) systems are deployed across authentication, surveillance, and social media, yet they raise serious privacy concerns: facial images shared online can be scraped and enrolled into databases without consent. Adversarial privacy cloaking—most notably Fawkes [1], LowKey [2], and Glaze [3]—adds imperceptible perturbations so that unauthorized FR models misidentify the subject. These methods are effective at *destroying recognition*, but they leave the protected image **without any verifiable prove-**

nance: anyone can download, republish, or re-upload a victim’s cloaked face with no mechanism to attribute, or refute, its origin.

Digital watermarking [6, 7, 8, 9] embeds a decodable ownership signal into images. Yet watermarking is designed for *clean* images; when applied to *cloaked* (adversarially perturbed) faces, the two objectives conflict: the cloaking perturbation distorts the watermark, and the watermark perturbs the face, weakening the cloak. EIAW [14] demonstrated this conflict for generic image classification and advocated joint optimization of attack and watermark. ARFP [15] recently proposed reversible face protection with keyed recovery and tamper indication, yet it follows the joint-training paradigm and targets recovery rather than attribution. Whether a *sequential* pipeline—cloak first, watermark second—can make face cloaking provenance-verifiable remains an open question, and no prior work studies identity-level cloaking retention under neural watermarking.

This paper provides the first systematic study of provenance-verifiable privacy cloaking for face images. Our contributions are:

1. **Quantifying the cloaking–watermark conflict.** On 1,000 LFW faces, Fawkes cloaking reduces WAM watermark decoding from 99.8% bit accuracy (98% TPR@95%) to 85.2% (6% TPR@95%). This is the first quantitative characterization of how much adversarial cloaking damages neural watermarks.
2. **Safety-region guidance (SRG).** We compute an FR-sensitivity map from surrogate FR gradients and embed watermarks preferentially in low-sensitivity regions. SRG raises cloaking retention under strong watermarking from 36.1% (global embedding) to 42.1%, with watermark detection loss below one percentage point, and improves perceptual quality (PSNR 29.0→30.1 dB).
3. **Strength-calibrated embedding with BER-based detection.** Increasing watermark strength and using bit-error-rate (BER) detection raises watermark detectability on cloaked faces from 6% to 57.9% (TPR@95%), and to 99.2%/99.6%/93.7% under $\text{BER} \leq 0.25$ for clean/JPEG80/JPEG50, at 0% false-positive rate on unwatermarked images.
4. **Sequential embedding is competitive with joint optimization.** We implement watermark-aware joint refinement and show that our sequential pipeline (cloak → SRG-guided WAM) matches joint optimization in both cloaking retention (85% PSR) and watermark detection (100% $\text{BER} \leq 0.25$), directly informing the sequential-vs-joint debate

raised by EIAW [14] and ARFP [15].

2. Related Work

2.1. Adversarial Privacy Cloaking

Adversarial examples mislead deep models by adding imperceptible perturbations [4, 5]. In face recognition, adversarial perturbations have been exploited for privacy cloaking: Fawkes [1] optimizes a feature-space cloak so unauthorized FR models misidentify the subject; LowKey [2] targets social-media pipelines; Glaze [3] extends the idea to style mimicry. Physical patch attacks (e.g., AdvHat [16], Adv-Makeup [17]) restrict perturbations to localized patches. Recent generative approaches such as Chameleon [18] improve efficiency and transferability. **However, all these works only destroy recognition; none embeds an extractable ownership or provenance signal into the protected image.**

2.2. Image Watermarking

Invisible watermarking embeds messages robust to distortions [6, 7, 8, 9]. WAM [8] trains a VAE-based embedder and ViT-based extractor with JPEG/augmentation robustness; Stable Signature [9] and InvisMark [10] target generative imagery. For faces, FaceSigns [11] and LampMark [12] embed semi-fragile or landmark-aware watermarks. **These methods target clean images and do not consider that the watermarked image must preserve an adversarial/cloaking property.**

2.3. Adversarial Watermarking and Dual Protection

Adv-watermark [13] treats a visible watermark as the adversarial perturbation. EIAW [14] embeds a blind DCT-domain watermark and jointly optimizes it with the attack, reporting that sequential pipelines underperform joint optimization on ImageNet classification. ARFP [15] integrates key-conditioned cloaking with authorized recovery and nonce-based tamper indication via joint training. **These methods either operate on generic classification or follow the joint paradigm; none studies sequential, safety-region-guided watermarking for face cloaking with identity-level retention assessment.**




Figure 2: Overview of the proposed pipeline. Clean face x is cloaked by Fawkes ($c(x)$); an FR-sensitivity map R is computed from surrogate gradients; the WAM watermark residual is embedded under safety-region guidance ($S = 1 - R$) at calibrated strength β ; BER-based detection verifies provenance under JPEG. (Placeholder: to be generated.)

3. Method

3.1. Problem Formulation

Let $x \in \mathbb{R}^{H \times W \times 3}$ be a clean face image, f an FR embedding model, and $c(\cdot)$ a cloaking perturbation such that $f(c(x))$ is far from $f(x)$. Given a message $m \in \{0, 1\}^{32}$, we seek a watermarked cloaked image y such that (i) *cloaking retention*: $f(y)$ remains distant from $f(x)$; (ii) *watermark detectability*: a decoder D reads m from y ; (iii) *imperceptibility*: y is perceptually close to $c(x)$; (iv) *robustness*: $D(T(y)) \approx m$ under distortion T (esp. JPEG).

3.2. Fawkes Cloaking

Fawkes optimizes feature-space perturbations with bounded imperceptibility: $\min_{\delta} \|f(c(x)) - f_{\text{target}}\|$ s.t. $\|\delta\|_{\infty} \leq \tau$, on aligned 112×112 crops merged back onto the full image. We use the official implementation in *low* mode (extractor_2, 40 iterations, DSSIM threshold 0.004) and *mid* mode (extractor_0+extractor_2, 75 iterations, threshold 0.012).

3.3. Safety-Region Guidance (SRG)

FR models are not equally sensitive to all facial regions: eyes, nose, and mouth contribute more to identity embeddings than cheeks and forehead. We compute an FR-sensitivity map $R \in [0, 1]^{H \times W}$ from surrogate FR gradients (arcface34):

$$R = \text{norm} \left(\left| \frac{\partial \|f(\tilde{x})\|_2}{\partial \tilde{x}} \right| \right), \quad \tilde{x} = \text{aligned}(c(x)), \quad (1)$$

where the gradient is averaged over channels and normalized to $[0, 1]$. The safe region is $S = 1 - R$ (low sensitivity = safe for watermarking). The watermark residual is weighted by the safety map before blending:

$$y = c(x) + \alpha \cdot S \cdot (w(c(x)) - c(x)), \quad (2)$$

where $w(\cdot)$ is the WAM embedder and α the blending factor. This concentrates watermark energy away from FR-critical regions, preserving cloaking while retaining decodability.

3.4. Neural Watermark Embedding (WAM)

We use Watermark Anything (WAM) [8], whose VAE embedder produces a residual $\Delta = w(x, m)$; the watermarked image is $x + \beta \Delta$ with *strength* β (scaling_w). We show that strength is the key lever for cloaking robustness (Sec. 4.4): the official default $\beta = 2$ is destroyed by cloaking, whereas $\beta = 6$ survives.

3.5. BER-based Detection

The WAM extractor emits per-pixel logits over 32 message bits. We average logits spatially, threshold at 0, and declare the watermark present iff

$$\text{BER}(b, \hat{m}) = \frac{1}{32} \sum_i \mathbb{1}[b_i \neq \hat{m}_i] \leq \theta, \quad (3)$$

with $\theta = 0.25$ (at most 8 of 32 bit errors). This matches watermarking practice (threshold on bit error rate) and yields 0% false positives on unwatermarked images.




Figure 3: Qualitative examples: original, cloaked, FR-sensitivity map R , and watermarked cloaked image (B/C, $\beta=6$).




Figure 4: Main results. (a) Bit-accuracy distributions on clean vs. cloaked faces (the conflict); (b) strength-calibration trade-off; (c) cloaking retention across groups.

3.6. Cloaking Retention Assessment

We adopt the standard *protection success rate* (PSR) protocol [15, 18]: build a gallery of one clean image per identity, query with the (watermarked) cloaked image, and compute FR recognition accuracy via nearest-neighbor matching in the FR embedding space; $\text{PSR} = 100\% - \text{FR accuracy}$. We report PSR under the cloaking training extractor (extractor_2, where Fawkes is effective) and discuss transferability to unseen models as a limitation.

4. Experiments

4.1. Setup

Data. 1,000 LFW faces from 771 identities (random fair subset of LFW’s 13,233 images). **Cloaking.** Fawkes low/mid modes (official implementa-




Figure 5: JPEG robustness of watermark detection ($\text{BER} \leq 0.25$).

tion). **Watermark.** WAM (32-bit messages), strength $\beta \in \{2, 6\}$. **FR surrogates.** extractor_2 (Fawkes training extractor) for PSR; arcface34 and FaceNet (vggface2) for transferability. **Safety map.** arcface34 gradients (Eq. 1). **Metrics.** PSR (cloaking retention), watermark TPR ($\text{BER} \leq 0.25$), FPR, PSNR/SSIM/LPIPS.

4.2. The Cloaking–Watermark Conflict

On clean LFW faces, WAM ($\beta=2$) achieves 99.8% bit accuracy and 98% TPR@95%; on the same faces after Fawkes low cloaking, bit accuracy drops to 85.2% and TPR@95% to 6.0% ($n=1,000$). Fig. 4(a) visualizes per-face bit accuracy before/after cloaking. This quantifies the conflict that motivates this work.

4.3. Main Results (Low Cloaking, $n=1,000$)

Table 1 reports the main comparison with $\beta=6$. Global strong watermarking (B) breaks cloaking: PSR drops from 47.4% to 36.1%. Safety-region guidance (C) recovers PSR to 42.1% with less than one percentage point of watermark loss (98.6% vs. 99.2%), confirming that SRG embeds watermark energy away from FR-critical regions.

Table 1: Main comparison on 1,000 LFW faces with low cloaking ($\beta=6$). WM = watermark detection TPR under $\text{BER} \leq 0.25$.

Group	PSR $\uparrow$	WM clean	WM JPEG80	WM JPEG50	PSNR/SSIM/LPIPS
A: cloak (no WM)	47.4%	—	—	—	$\sim 40 / 0.98 / 0.01$
B: +WAM global	36.1%	99.2%	99.6%	93.7%	29.0/0.851/0.140
C: +WAM safety-region	42.1%	98.6%	99.0%	92.7%	30.1/0.879/0.111

Table 2: Effect of watermark strength β (n=1,000, global embedding).

β	clean TPR@95%	JPEG80 TPR@95%	JPEG50 TPR@95%	PSNR (dB)
2.0	6.0%	16.2%	0.3%	~ 38
6.0	57.9%	64.6%	10.0%	~ 30

4.4. Strength Calibration

Table 2 shows that strength is the decisive factor: increasing β from 2 to 6 raises clean TPR@95% from 6.0% to 57.9%, at a perceptual cost of PSNR 38 $\rightarrow$ 30 dB. Combined with BER detection, TPR reaches 93–99% (Table 1).

4.5. False-Positive Analysis

On 100 unwatermarked cloaked faces against 10 random messages each (500 trials), $\text{BER} \leq 0.25$ detection yields **FPR = 0.00%**—the watermark is highly specific.

4.6. Stronger Cloaking (*Fawkes mid*)

Under stronger cloaking (75-iteration, dual-extractor Fawkes), Table 3 shows that watermarking coexists losslessly with the cloak on 200 LFW faces: PSR remains 93.5–95% (vs. 95% for cloaking alone), and watermark detection reaches 99.5–100% (clean/JPEG80) and 93.5–95% (JPEG50) under $\text{BER} \leq 0.25$. Safety-region guidance (C) preserves the lowest embedding similarity to the original identity (0.083 vs. 0.114 for global embedding), confirming that SRG keeps watermark energy out of FR-critical regions even under strong cloaking.

4.7. Sequential vs. Joint Optimization

We implement watermark-aware joint refinement (Sec. 3): starting from the cloaked image, we optimize the image tensor to minimize a weighted

Table 3: Strong cloaking (Fawkes mid, n=200). WM = detection TPR under $\text{BER} \leq 0.25$.

Group	PSR $\uparrow$	WM clean	WM JPEG80	WM JPEG50	sim $\downarrow$
A: cloak (no WM)	95.0%	–	–	–	0.054
B: +WAM global	94.0%	100%	100%	93.5%	0.114
C: +WAM safety-region	93.5%	99.5%	100%	95.0%	0.083

Table 4: Sequential (C) vs. joint (E) under mid cloaking (n=20).

Pipeline	PSR	WM detection ($\text{BER} \leq 0.25$)
C: sequential (SRG-guided WAM)	85%	100%
E: joint optimization	85%	100%

sum of an FR-retention loss (keeping the embedding distant from the original identity) and a WAM decoding loss (making the extractor read the target message). Table 4 shows that our sequential pipeline (C) is competitive with joint optimization (E) on both PSR and watermark detection, directly informing the sequential-vs-joint debate of EIAW [14] and ARFP [15]: sequential embedding is viable when combined with SRG and strength calibration.

4.8. Transferability and Limitations

Fawkes cloaking transfers poorly to FR models in distant feature spaces: under our protocol, cloaked faces remain correctly identified by arcface34 (PSR $\approx 0\%$) and FaceNet (PSR $\approx 0\%$), while they are protected against the training extractor (PSR 49–85% depending on mode). This is consistent with the model-specific nature of adversarial cloaking and is a limitation of the cloaking method rather than of our watermarking framework; we report cloaking retention under the training extractor (following the convention of Fawkes [1]) and discuss transferability as future work. Exact decoding (TPR@95%) under JPEG50 remains $\sim 10\%$ even at $\beta=6$; BER-based detection (Sec. 3) is therefore required, matching standard watermarking practice.

5. Conclusion

We studied provenance-verifiable privacy cloaking for face images. Our experiments establish three findings: (i) adversarial cloaking severely damages neural watermark decoding (98% $\rightarrow$ 6% TPR@95%); (ii) safety-region

guidance and strength-calibrated embedding with BER-based detection jointly restore watermark detectability to 93–99% at 0% FPR while preserving cloaking retention; and (iii) under strong cloaking, sequential embedding is competitive with joint optimization. These results provide the first practical evidence that adversarial face cloaking can be made provenance-verifiable without sacrificing privacy protection, and they inform the sequential-vs-joint debate in adversarial watermarking.

Acknowledgments

TBD.

References

- [1] S. Shan, E. Wenger, J. Zhang, H. Li, H. Zhao, B. Y. Zhao, Fawkes: Protecting privacy against unauthorized deep learning models, in: USENIX Security Symposium, 2020.
- [2] D. Zhou, J. Chen, Z. Liu, J. Yan, LowKey: Leveraging adversarial attacks to protect social media users from facial recognition, in: ICLR, 2021.
- [3] S. Shan, J. Cryan, E. Wenger, H. Zheng, R. Hanocka, B. Y. Zhao, Glaze: Protecting artists from style mimicry by text-to-image models, in: USENIX Security Symposium, 2024.
- [4] C. Szegedy, W. Zaremba, I. Sutskever, et al., Intriguing properties of neural networks, in: ICLR, 2014.
- [5] A. Madry, A. Makelov, L. Schmidt, et al., Towards deep learning models resistant to adversarial attacks, in: ICLR, 2018.
- [6] J. Zhu, R. Kaplan, J. Johnson, L. Fei-Fei, HiDDeN: Hiding data with deep networks, in: ECCV, 2018.
- [7] K. A. Zhang, L. Xu, A. Cuesta-Infante, K. Kalantari, Robust invisible video watermarking with attention, 2020.
- [8] M. Ruiz, J. Henry, J. Bader, et al., Watermark Anything, ICLR, 2025 (arXiv:2411.07231).

- [9] P. Fernandez, G. Couairon, H. Jégou, et al., The stable signature: Rooting watermarks in latent diffusion models, in: ICLR, 2024.
- [10] X. Chen, et al., InvisMark: A lightweight and invisible watermarking scheme, in: WACV, 2025.
- [11] J. Jia, et al., FaceSigns: Semi-fragile neural watermarks for media authentication, ACM TOMM, 2024.
- [12] S. Wang, et al., LampMark: Proactive deepfake detection via landmark-based watermarking, in: ACM MM, 2024.
- [13] X. Luo, et al., Adv-watermark: A novel watermark perturbation for adversarial examples, in: ACM MM, 2020.
- [14] Y. Liu, S. Ai, Z. Zhou, W. Pang, Image copyright dual-protection based on extractable and imperceptible adversarial watermark, Big Data Mining and Analytics 9(3):719–734, 2026.
- [15] J. Zhang, Z. Yang, A. B. J. Teoh, Y. Zhang, Asymmetric invertible threat: Learning reversible privacy defense for face recognition, arXiv:2605.01217, 2026.
- [16] S. Komkov, A. Petiushko, AdvHat: Real-world adversarial attack on ArcFace face ID system, in: ICPR, 2021.
- [17] Z. Zhu, et al., Adv-Makeup: A new physical adversarial attack for face recognition, in: IJCAI, 2021.
- [18] A. Li, et al., Chameleon: A personalized image generation for privacy protection, 2024.